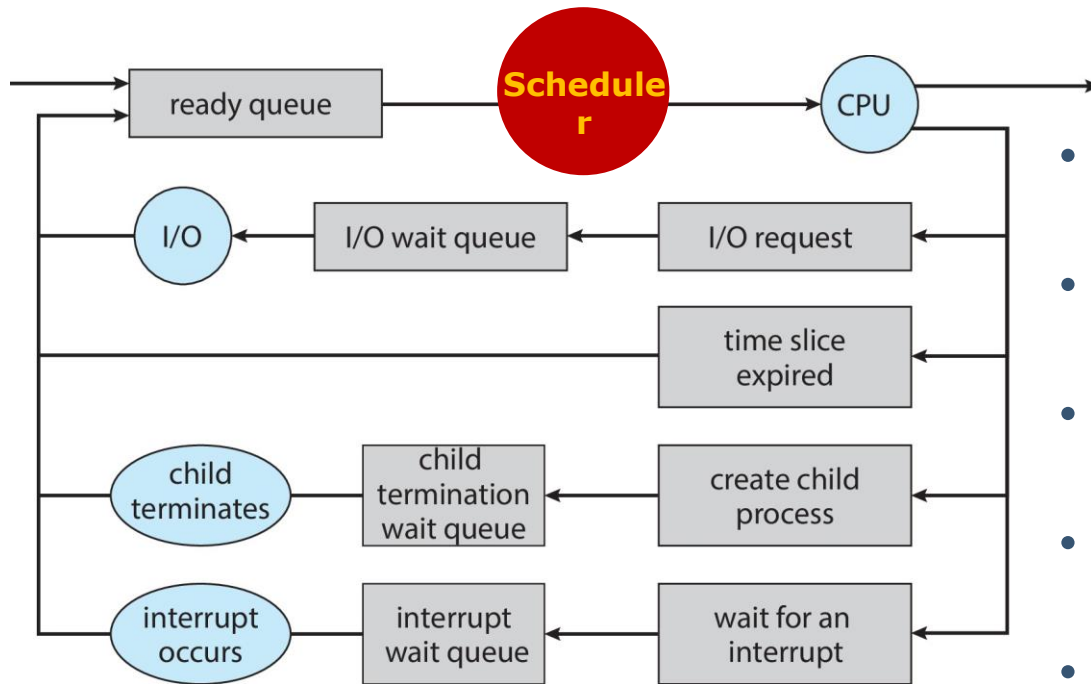


Process Management

- Scheduling Policies-I -

Recall: Basics and Criteria



- Maximizing CPU Utilization
- Maximizing Throughput
- Minimizing Turnaround Time
- Minimizing Waiting Time
- Minimizing Response Time

First Come First Serve (FCFS)

- Schedule process that requested execution first
 - Oldest process in the ready queue
 - Manage ready queue using a first-in, first-out policy.
 - When CPU idle, scheduler selects head of line (HOL)
 - Once scheduled, process is removed from the queue, next one become HOL
 - New process added to queue tail



FCFS Example

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	-3	24
P_2	3	3
P_3	9	3

The Gantt Chart for the schedule is:



Process	Waiting Time
P_1	$0 - (-3) = 3$ ms
P_2	$24 - 3 = 21$ ms
P_3	$27 - 9 = 18$ ms
Average	14 ms

FCFS Scheduling (Cont.)

Suppose that the processes arrive in the order:

$$P_2, P_3, P_1$$

- The Gantt chart for the schedule is:



- Waiting time for $P_1 = 6$; $P_2 = 0$; $P_3 = 3$
- Average waiting time: $(6 + 0 + 3)/3 = 3$
- Much better than previous case
- **Convoy effect** - short process behind long process
 - Consider one CPU-bound and many I/O-bound processes
 - Is there a convoy effect in the above Gantt chart?

FCFP Pros and Cons

- Pros:
 - Simple and easy to implement
 - Fair – every process will eventually run as long as CPU is not blocked by a process
- Cons:
 - Waiting time depends on arrival order
 - Short process stuck waiting for longer process – convoy effect
 - No differentiation between processes
 - Nonpreemptive


Shortest Job First (SJF)

- Schedule process with shortest burst time first
 - Associate with each process the length of its next CPU burst ?
 - Use these lengths to schedule the process with the shortest time
 - Waiting-time optimal, but poor on turnaround and response times
- Example (nonpreemptive SJF):

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	-4	6
P_2	-3	8
P_3	-2	7
P_4	-1	3

- Average waiting time FCFS = $(4 + 9 + 16 + 22)/4 = 12.75$ ms
- Average waiting time SJF = $(1 + 7 + 11 + 19) / 4 = 9.5$ ms

- SJF




Predicting Burst Times

- Predicting ?

- Knowing exact burst times is impossible, must be predicted
- Typically, should have some correlation with previous bursts

- Exponential average based prediction:

- τ_n : Predicted value for n -th burst
- t_n = Actual time of n -th burst
- $\tau_{n+1} = \alpha \cdot t_n + (1 - \alpha) \cdot \tau_n$
- α must be in $[0,1]$, typically $= 1/2$



Preemptive SJF

- Shortest Remaining Time First (SRTF)
 - Stop current process and schedule new one if burst time for latter shorter than remaining time former.
 - Now we add the concepts of varying arrival times and preemption to the analysis considering four processes and the length of CPU burst in milliseconds (msec).


<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

Waiting Time in SRTF

- **Waiting time** = Start time of last executed segment – Arrival time – Scheduled time of all prior segments
- Example

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

- Waiting time (P_1) = $10 - 0 - 1 = 9$ ms
- Waiting time (P_2) = $1 - 1 - 0 = 0$ ms
- Waiting time (P_3) = $17 - 2 - 0 = 15$ ms
- Waiting time (P_4) = $5 - 3 - 0 = 2$ ms



- Average waiting time = $(9+0+15+2)/4 = 26/4 = 6.5$ ms

Round Robin (RR)


- FCFS with preemption to enable time sharing
 - Defines a time quantum/slice (10 - 100 ms)
 - **Circular operation**
 - Process bursts sequentially get up to 1 slice
 - First burst get another slice after slice of last one expires, and so on
- Implementable in a circular FIFO queue
 - Timer interrupts every quantum to schedule next process
 - Interrupt occurs = context switching, current process sent to tail
 - A terminating process releases CPU for next process
- Pros: Fast response time
- Con: Long waiting and turnaround times

Example

The process will
complete at
Round 5

Time slice: 4 ms

Process	Arrival Time	Burst Time	After R1	After R2	After R3	After R4
P_1	0	14	10	6	2	-
P_2	1	11	7	3	-	-
P_3	9	15	-	11	7	3




Process	Waiting Time	Turnaround Time	Response Time
P_1	$31 - 0 - 12 = 19$ ms	$33 - 0 = 33$ ms	$0 - 0 = 0$ ms
P_2	?	?	$4 - 1 = 3$ ms
P_3	?	?	$16 - 9 = 7$ ms
Average	?	?	3.33 ms

RR Facts and Considerations

- N bursts in ready queue, t_s ms per slice
 - Each burst gets $\leq 1/N$ of the time
 - No burst waits more than $(N - 1) \times t_s$ without executing
 - Response time $\leq (N - 1) \times t_s$
- t_s must be carefully selected
 - $t_s \gg \Rightarrow$ FCFS
 - $t_s \ll \Rightarrow$ Too many context switching \Rightarrow Low CPU utilization
 - $t_s \ll \Rightarrow$ Typically higher average turnaround and waiting times

80% of CPU burst times $\leq t_s$

Time Quantum and Context Switch Time



Assume, a process is of 10 time unit:

- If $q = 12$ time unit, process finishes in 1 time q , no context switch
- If $q = 6$ time unit, process finishes in 2 time q , 1 context switch
- If $q = 1$ time unit, process finishes in 10 time q , 9 context switch