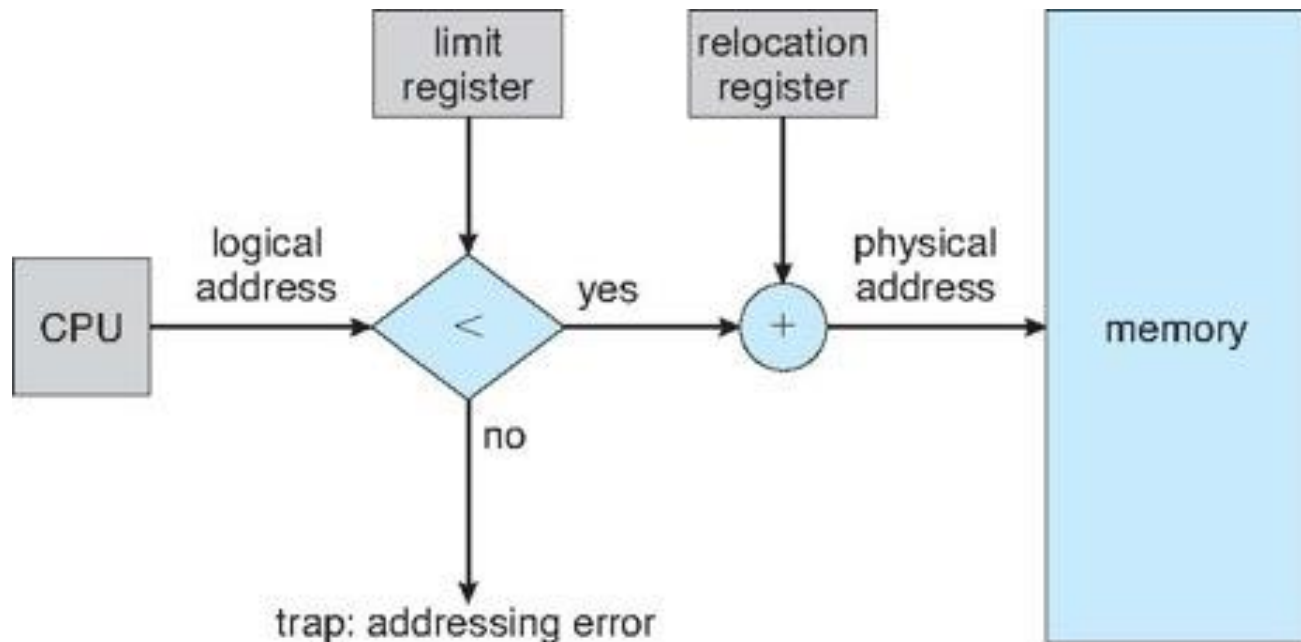


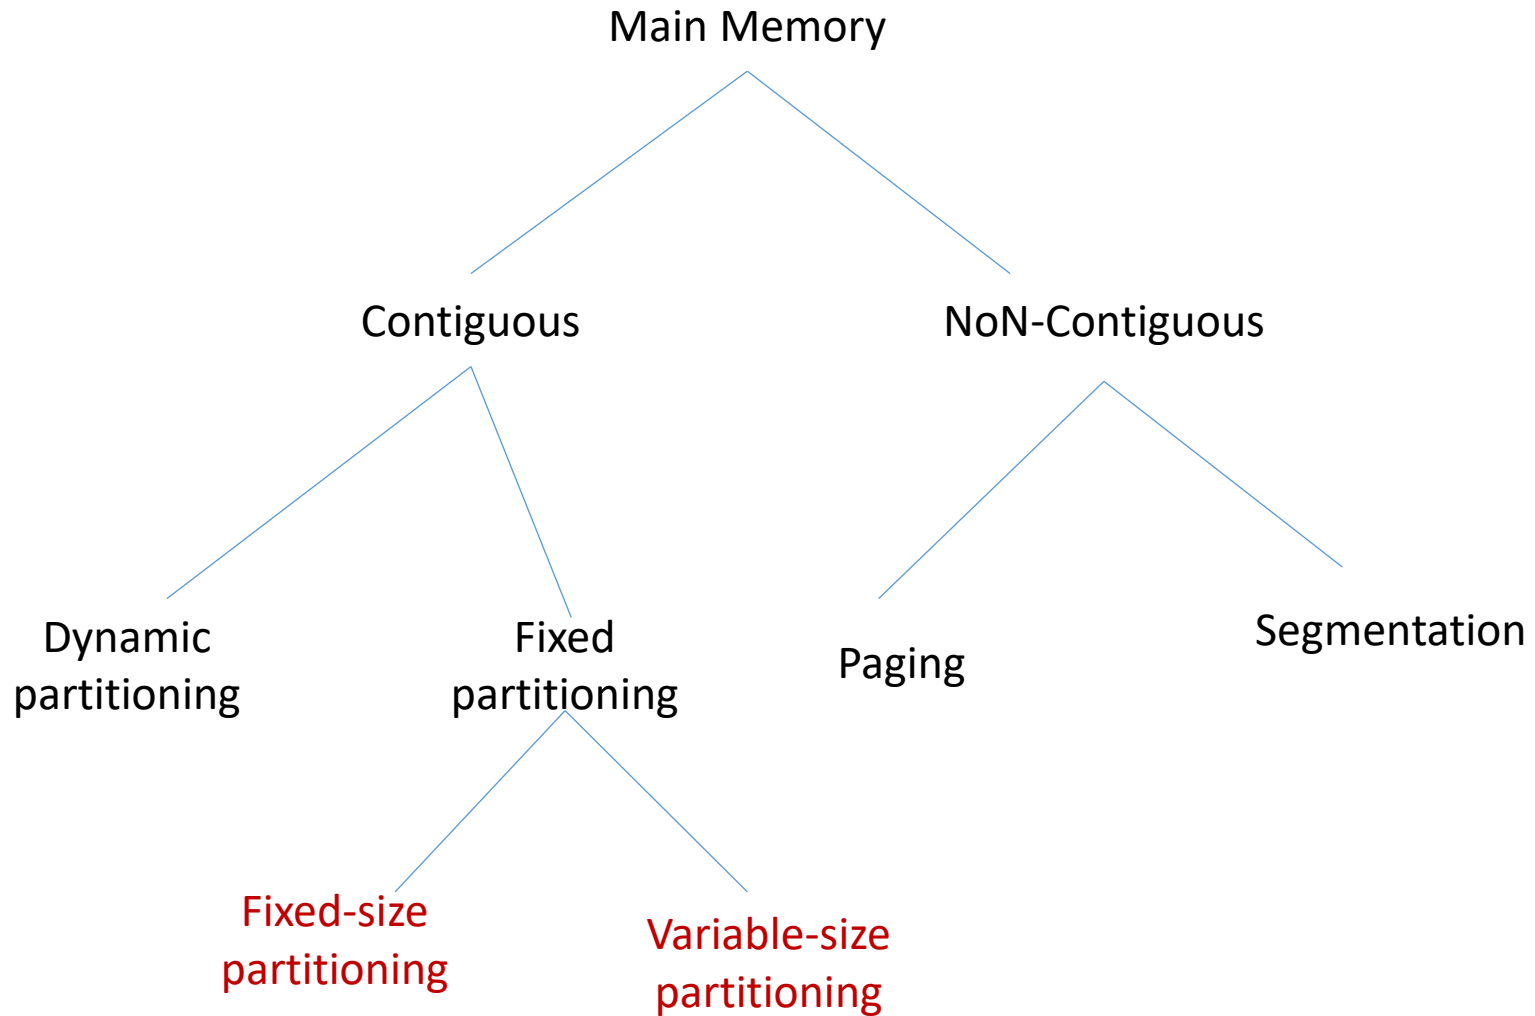
Chapter 9: Memory Management – Paging overview

Recall: Address Translation

- Logical addresses generated for each program
 - May have the same address space
 - MMU translates logical addresses into physical addresses
- Translation example



Recall: Memory Allocation




Paging

- Non-contiguous allocation approach
 - Allows parts of program to load in non-adjacent locations in memory
 - Resolves external fragmentation
 - Allows program parts to fill varying-sized memory holes (fragments)
- Pages: Fixed and equal-sized blocks of logical address space
 - Size typically in powers of 2; 4 KB (2^{10} bytes) to 4 MB (2^{22} bytes) per page
 - In Linux: type `getconf PAGESIZE`
- Frames: Fixed-sized blocks of the physical memory
- Paging
 - Loads pages of programs into any available frames in the memory
 - Program of size N pages loaded in N free frames
 - Page size is equal to frame size

Page Number and Offset


- Logical address division

- 2^m logical addresses divided into pages of size 2^n
- $m-n$ most significant bits of any address \Rightarrow Page number (p)
- n least significant bits of any address \Rightarrow Offset (d) inside this page



- Page number (p) – used as an index into a page table which contains base address of each page in physical memory
- Page offset (d) – combined with base address to define the physical memory address that is sent to the memory unit

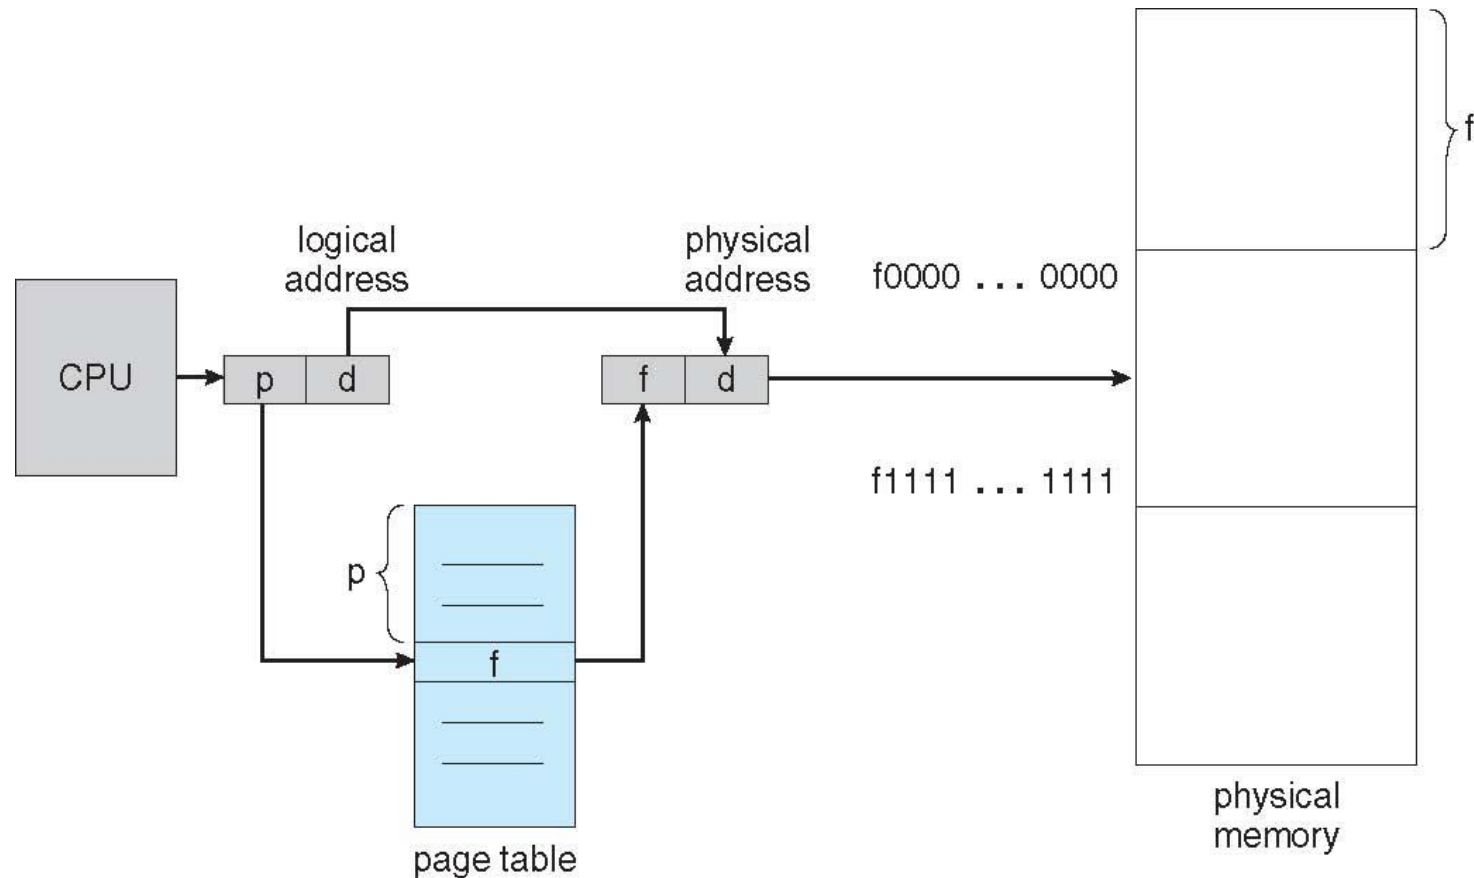
Page Number and Offset



- Example

- 64 kB logical space divided into pages of 1 kB $\Rightarrow m = 16, n = 10$
- **No. of pages** = $2^{(16-10)} = 2^6 = 64$ pages (pages 0 to 63)
- A logical address 0xA934 = **1010 1001 0011 0100** (A=10 decimal)
 - $p = 101010 \Rightarrow$ **First address on page** = **1010 1000 0000 0000** = 0xA800
 - $d = 0000 0001 0011 0100 = 0x0134$
 - Address = First address on page + d = 0xA800 + 0x0134 = 0A934

Address Translation in Paging- paging hardware



For each process, maintain 1 page table as each process will have page 0, page 1, and so on

Address Translation in Paging

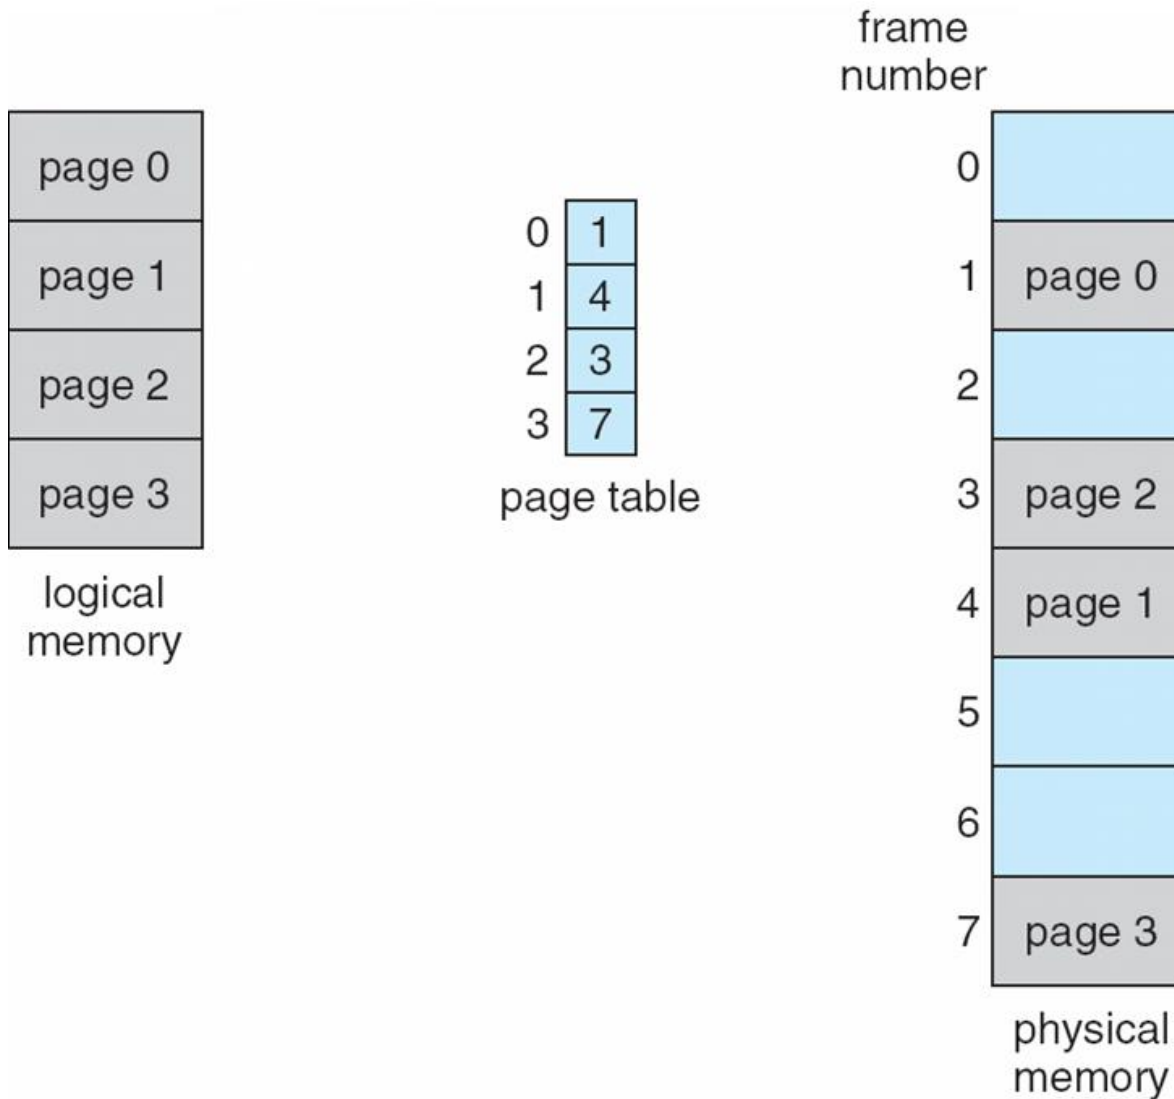
- **Page Table**

- Data structure associated to each process
- Page number \rightarrow Physical base address for frame
- Filled at process creation
- Offset of logical address points to physical address in frame starting from base address


- **Translation Steps**

1. Extract p and d
2. Get p 's frame base address (f) from table
3. Append/add d to f to obtain address in memory

Paging Model of Logical and Physical Memory

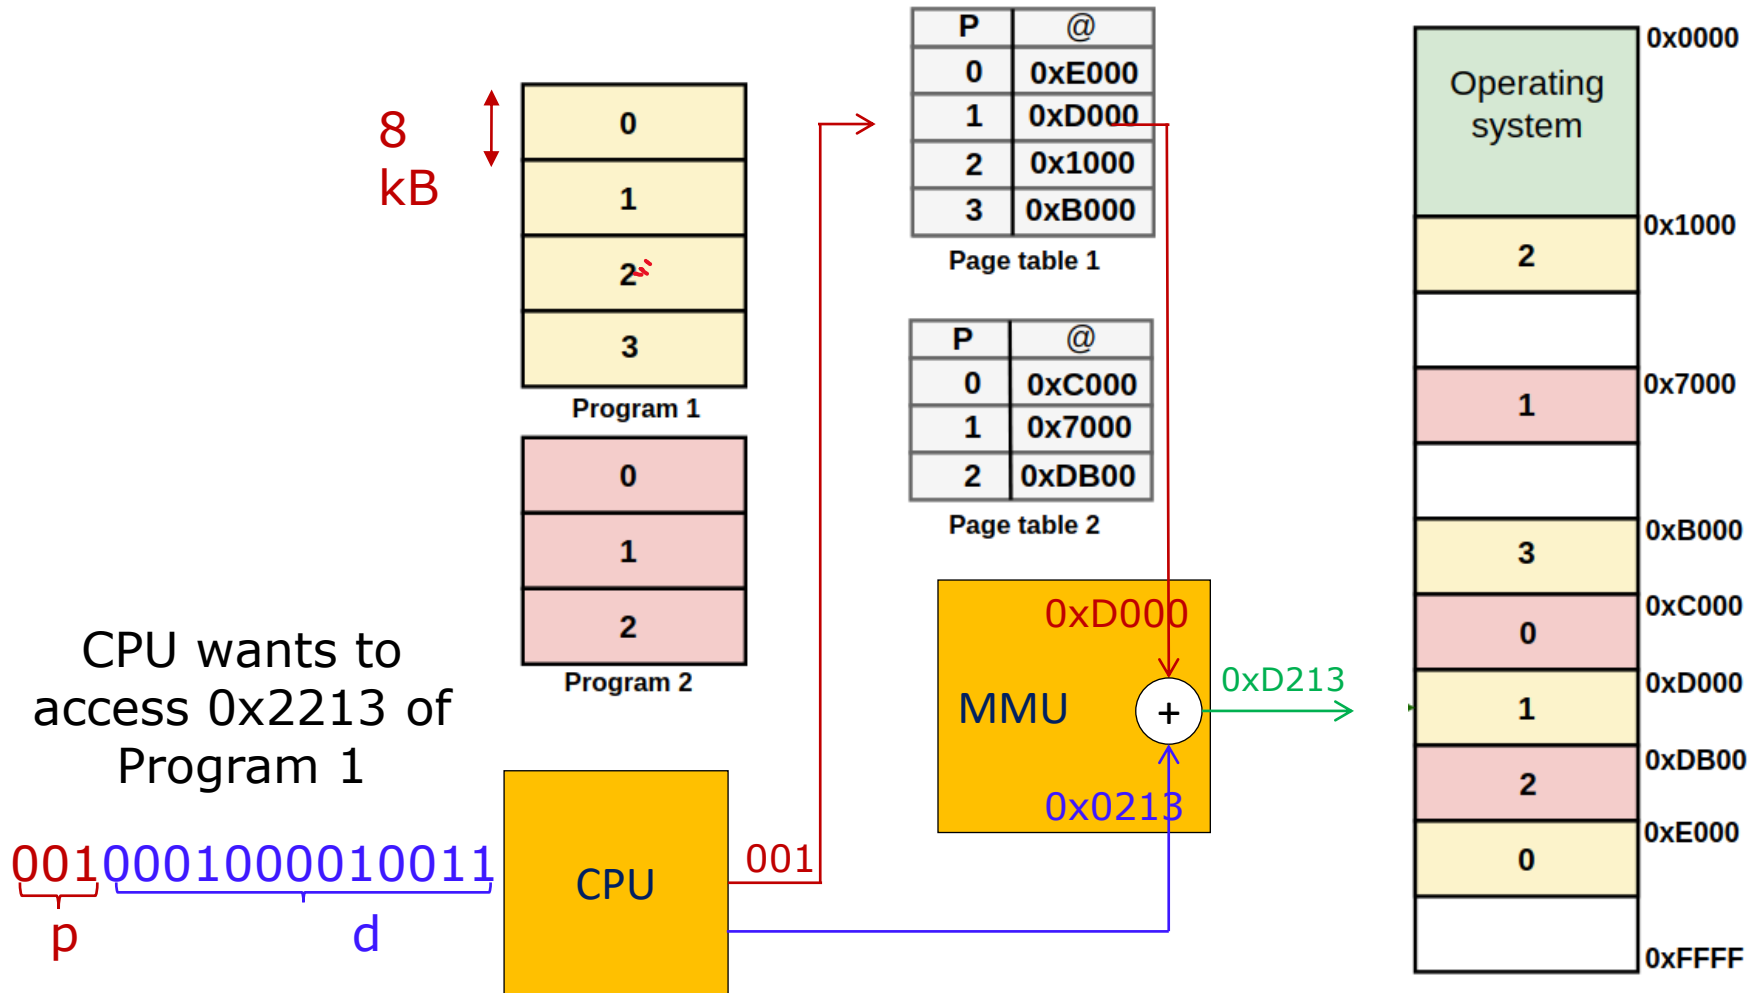


Paging Example



$n=2$ and $m=4$ 32-byte memory and 4-byte pages

Example



Internal Fragmentation

- May still occur with paging
 - Last frame not filled with information in last page of program
- Example:
 - Page/frame size = 2,048 bytes
 - Process size = 72,766 bytes
 - $72766/2048 = 35 \text{ pages} + 1,086 \text{ bytes}$
 - Internal fragmentation of $2,048 - 1,086 = 962 \text{ bytes}$
 - Allocation: 36 frames - one having 962 unused bytes
- Statistics
 - Worst case: process that is split into n pages plus 1 byte would be allocated $n + 1$ frames
 - Average case: $\frac{1}{2}$ page/frame size
- Solution: Smaller page/frame sizes
 - Larger overhead in keeping and searching larger page tables

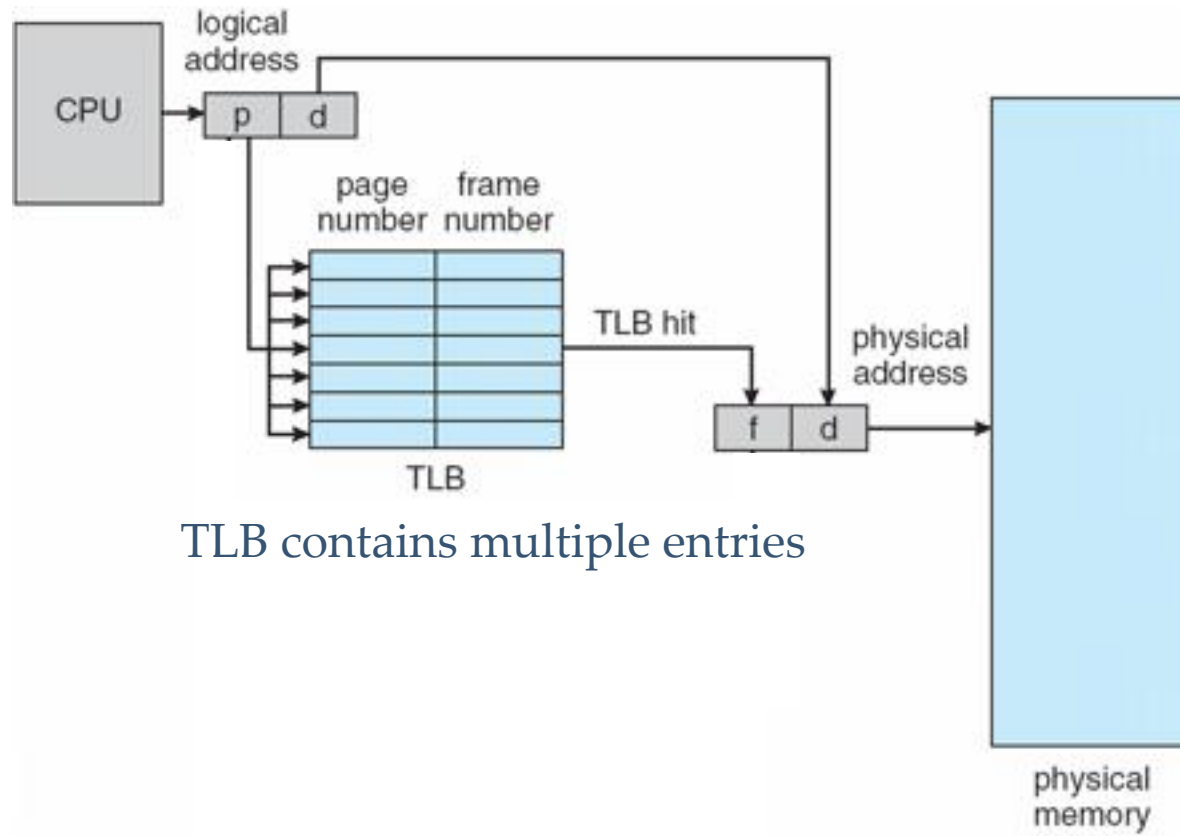
Implementation of Page Table

- Via Dedicated High-speed Registers
 - Each register storing one entry in the page table
 - **Pro:** Very fast in address translation
 - **Con:** Very slow context switching (content of each register must be exchanged by the page table entry of the next process..)
 - **Fact:** Satisfactory for small page tables up to 2^8 entries
- Via Memory (case of most contemporary computer)
 - Suitable for larger (most common) tables (e.g., 2^{20} entries)
 - Pointer to page table base address in memory stored in PCB
 - **Page-table base register (PTBR)**
 - Pointer loaded from PCB into PTBR at context switching
 - **Pro:** Much faster context switching time (changing processes require only change in this register)
 - **Con:** Slower memory access time – Two memory access to reach physical address: one to **locate the page table** and the other to access the identified **physical address**

Translation Look-Aside Buffer (TLB)

- Solution to 2-step memory access problem → special fast-lookup hardware cache called the translation look-aside buffer (or TLB)
 - Stores the recently used virtual-to-physical address translations to speed up the memory access process
- TLB Used for Page Search:
 - TLB stores few page-table entries between 32 and 1,024
 - Sometimes two TLBs used, one for instructions and one for data
 - MMU checks if page number is in TLB:
 - If so: **Hit** - Base address extracted, used to access memory
 - If not: **Miss** - 2-step memory access: get the address from the page table
- Replacement applied on TLB:
 - Add recently called address in TLB instead of an existing one
 - Replacement policies: LRU (least recently used), RR, Random

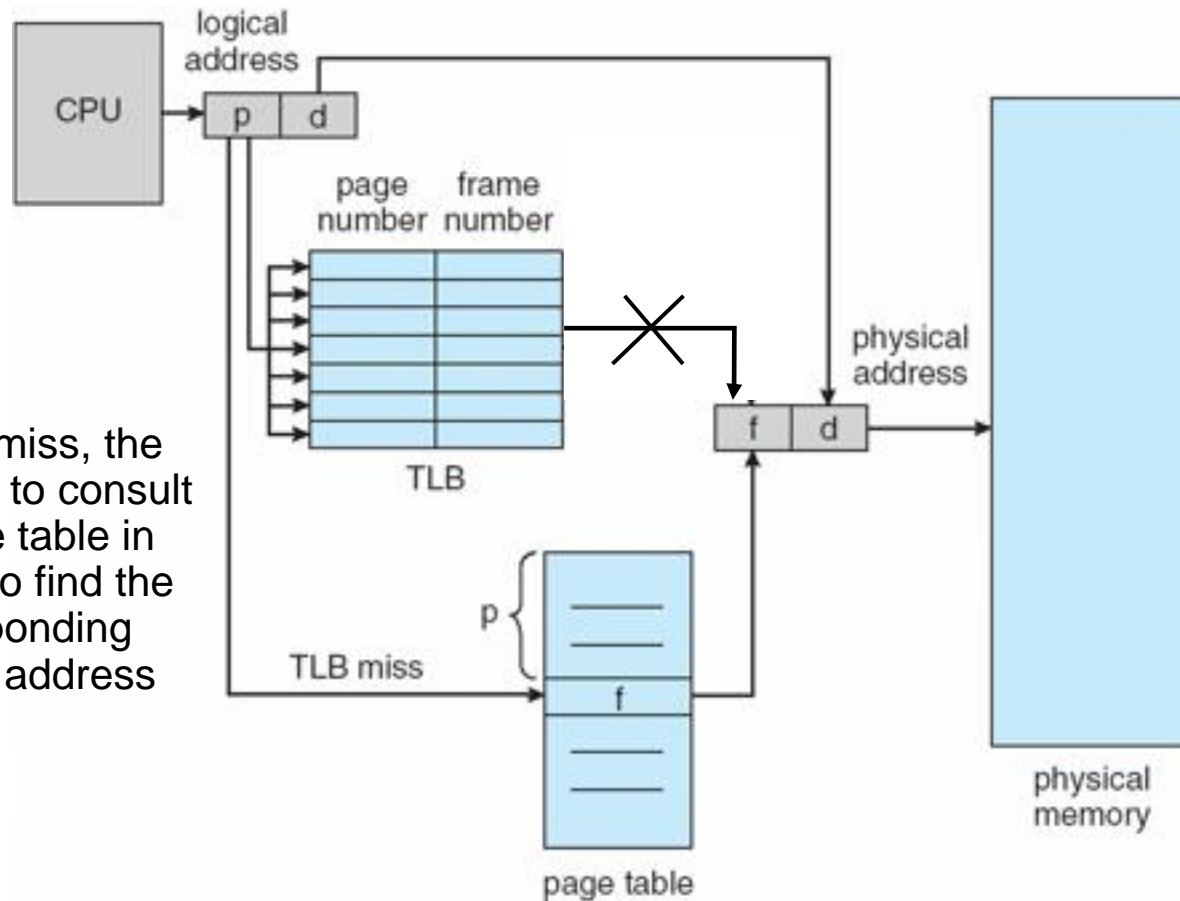
TLB-Hit



TLB contains multiple entries

TLB-Miss

On TBL-miss, the CPU need to consult the page table in memory to find the corresponding physical address



Effective Access Time (EAT)

- Access time differs between TLB hit and miss cases
 - EAT: Defines the overall performance of the system
- Calculation
 - **Hit Ratio:** Proportion of time a page is found in TLB
 - **Miss Ratio:** Proportion of time a page is not found ($1 - \text{Hit Ratio}$)
 - **Access Time:** Time needed to access the memory once

$$\text{EAT} = \text{Hit Ratio} \times \text{Access Time} + \text{Miss Ratio} \times 2 \times \text{Access Time}$$

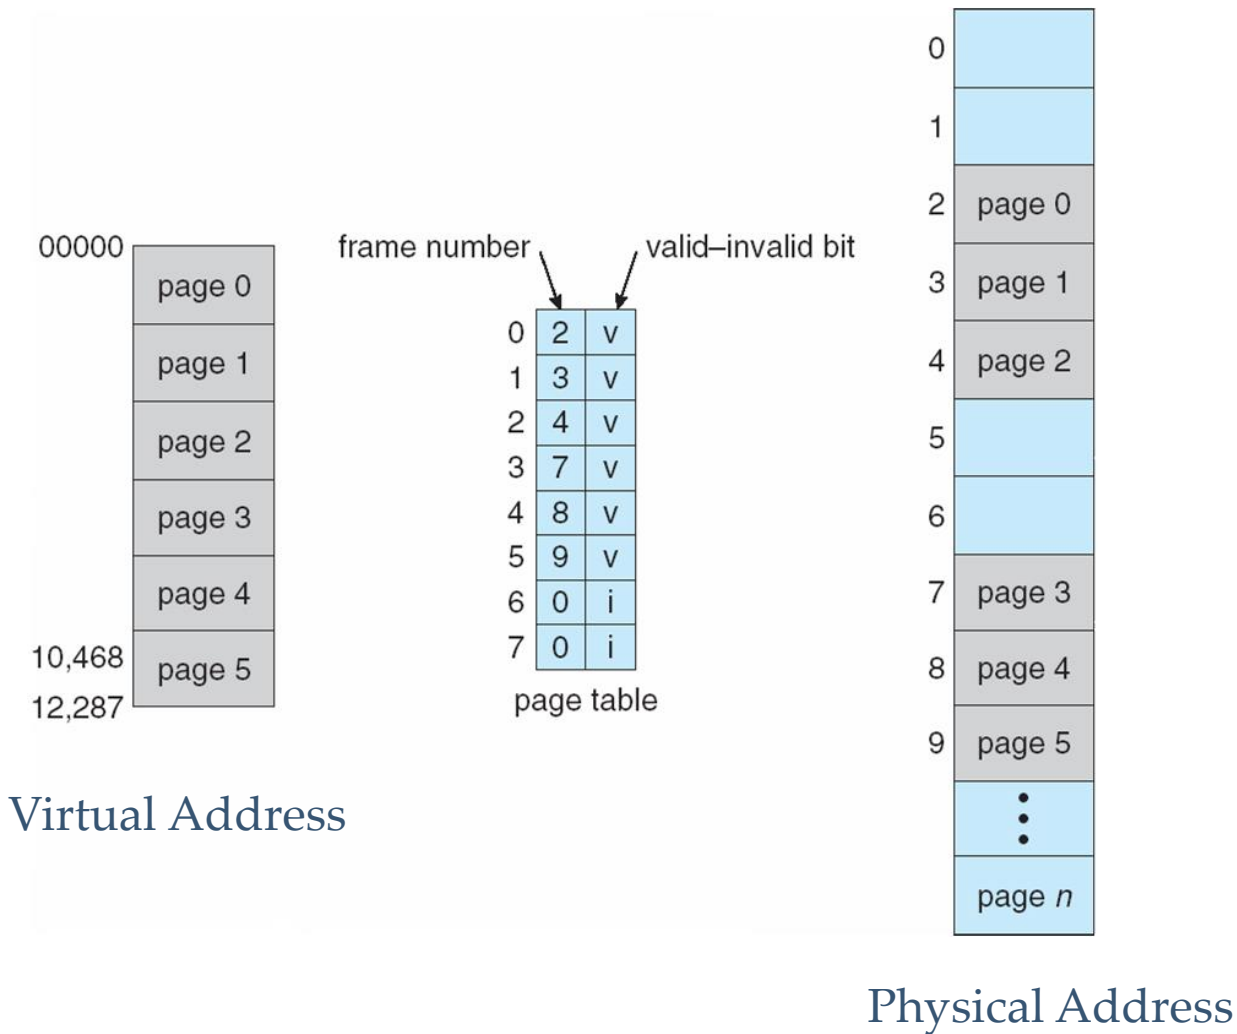
(2 is for 2-step memory access)

- Example
 - Hit Ratio = 0.8
 - Access Time: 2 ns
 - $\text{EAT} = 0.8 \times 2 + 0.2 \times 2 \times 2 = 2.4 \text{ ns}$

Memory Protection

- Memory protection implemented by associating protection bit with each frame to indicate if read-only or read-write access is allowed
 - Can also add more bits to indicate page execute-only, and so on
- Valid-invalid bit attached to each entry in the page table:
 - “valid” indicates that the associated page is in the process’ logical address space, and is thus a legal page
 - “invalid” indicates that the page is not in the process’ logical address space
 - Or use page-table length register (PTLR)
- Any violations result in a trap to the kernel

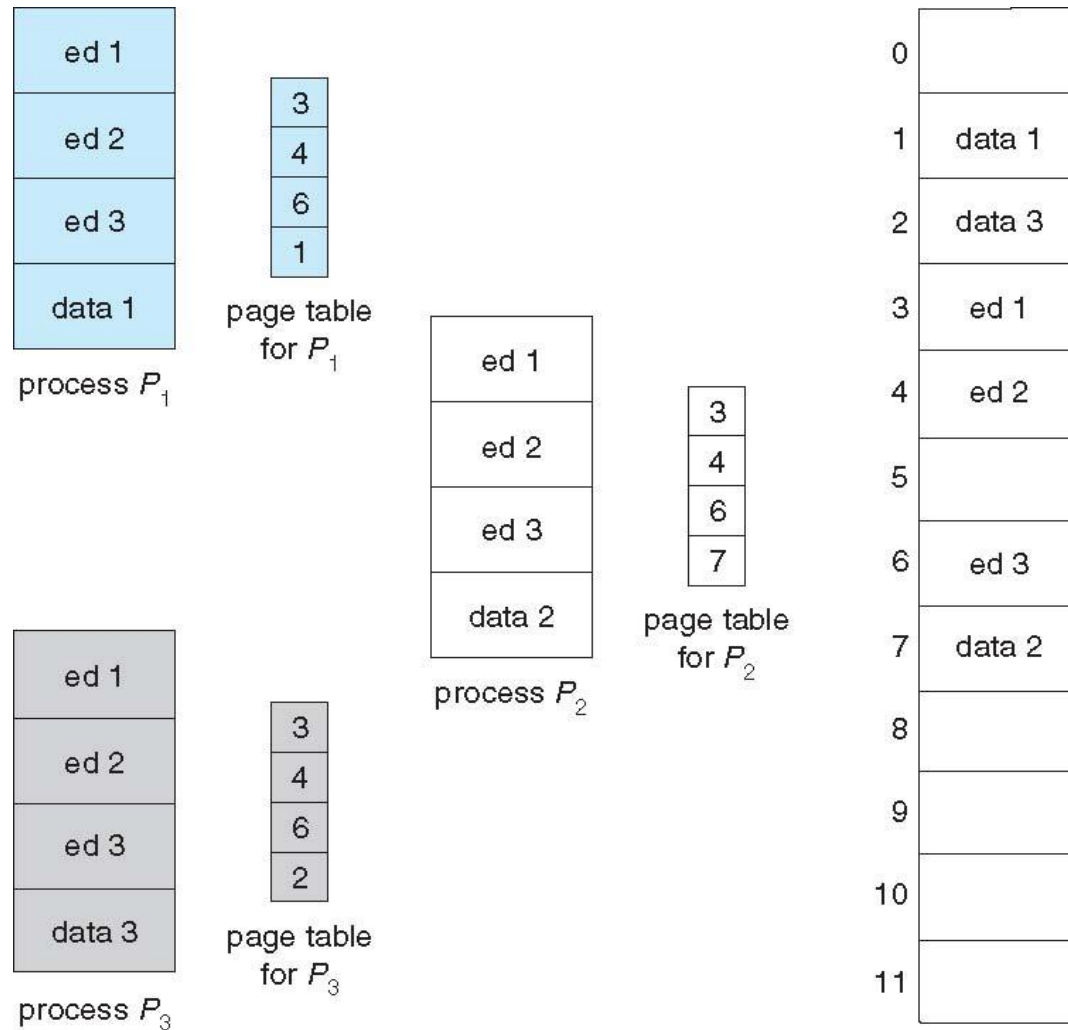
Valid (v) or Invalid (i) Bit In A Page Table



Shared Pages

- Shared code
 - One copy of read-only (reentrant) code shared among processes (i.e., text editors, compilers, window systems)
 - Similar to multiple threads sharing the same process space
 - Also useful for interprocess communication if sharing of read-write pages is allowed
- Private code and data
 - Each process keeps a separate copy of the code and data
 - The pages for the private code and data can appear anywhere in the logical address space

Shared Pages Example



Page Table Structures

- Modern computers support large logical address spaces
 - 32-bit systems: 2^{32} addresses
 - 64-bit systems: 2^{64} addresses.
- Excessively large page tables
 - Example: For 32-bit system
 - 4 kB page size (2^{12}) results in $2^{32}/2^{12} = 2^{20} \approx 1\text{M}$ table entries
 - If each entry is only 4 bytes, page table size = 4 MB
 - That amount of memory used to cost a lot
- Smarter Structures
 - Hierarchical Page Tables
 - Hashed Page Tables
 - Inverted Page Tables


Hierarchical Page Table

- Break up the logical address space into multiple page tables
- A simple technique is a two-level page table
- We then page the page table


Hierarchical Paging

- Dividing table into smaller pages (paging in page table)

- Remember: 1-level paging



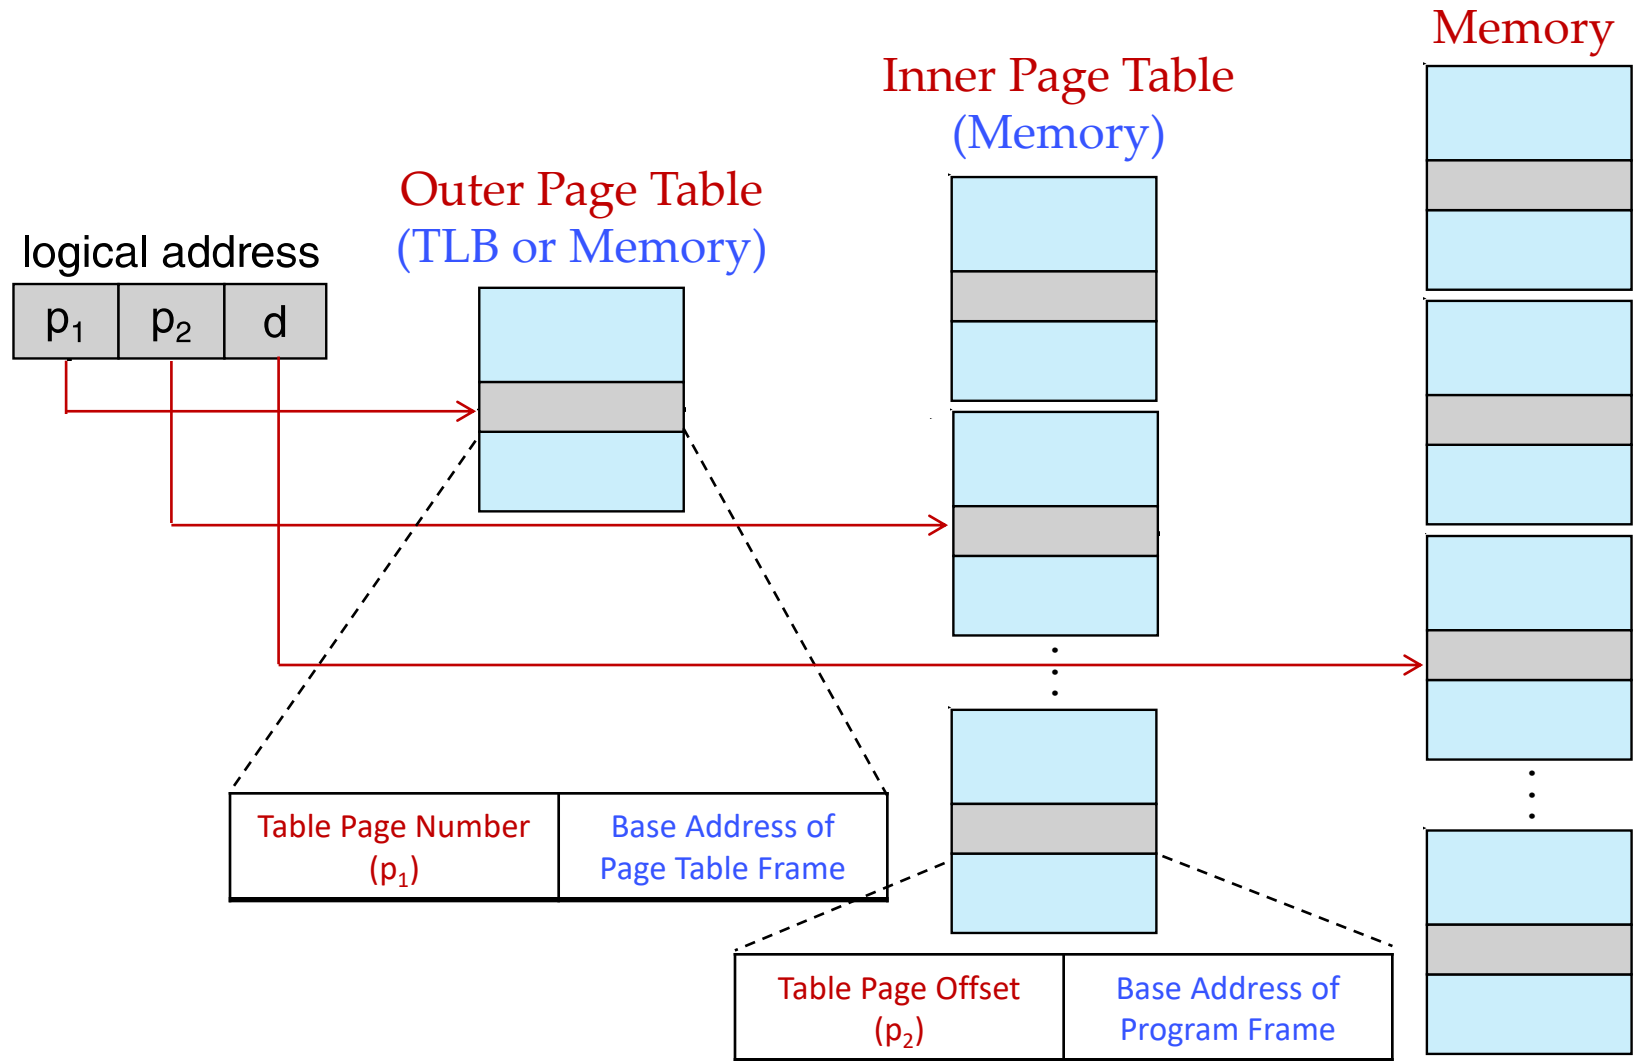
- Hierarchical (2-level) paging




- Two types of frames

- Page table frames:** Frames in the memory storing pages of the page table
- Program frames:** Frames in the memory storing pages of program

Address Translation




Two-Level Page-Table Scheme



Two-Level Paging Example

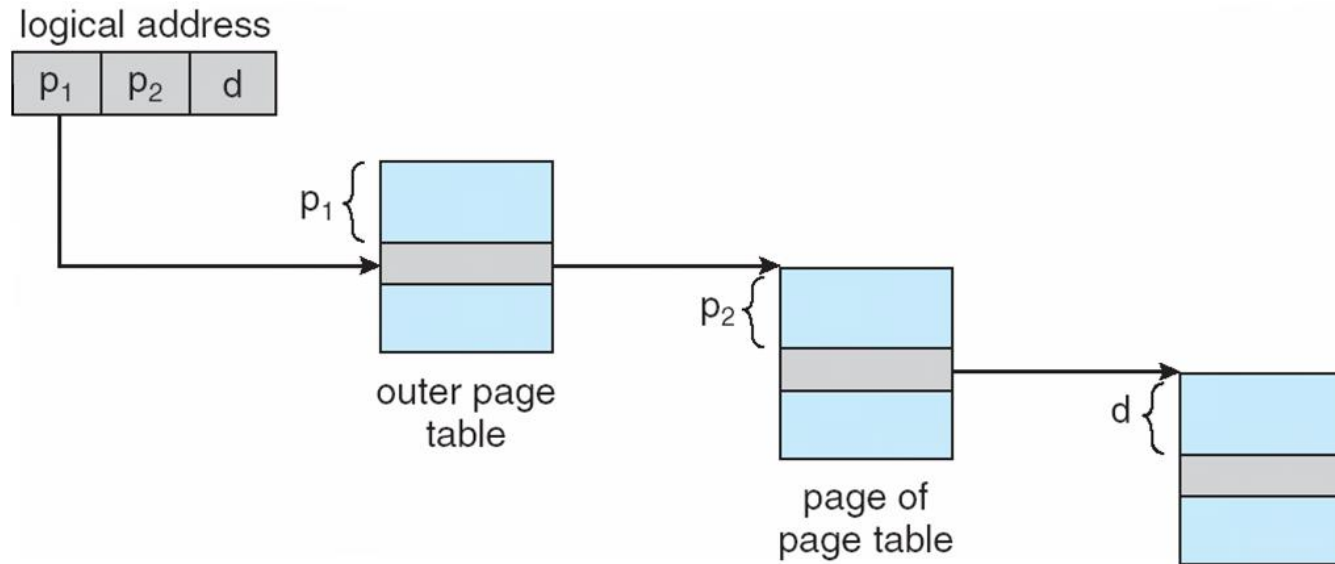
- A logical address (on 32-bit machine with 1K page size) is divided into:
 - a page number consisting of 22 bits
 - a page offset consisting of 10 bits
- Since the page table is paged, the page number is further divided into:
 - a 12-bit page number
 - a 10-bit page offset
- Thus, a logical address is as follows:



where p_1 is an index into the outer page table, and p_2 is the displacement within the page of the inner page table


- Known as forward-mapped page table

Address-Translation Scheme




Benefits

- Resolve huge page table problems
 - Storage: Table partitioned into 2^{m_1} pages stored in non-contiguous frames
 - Search: Drops down from $O(2^{(m_1 + m_2)})$ in the conventional 1-level paging to $O(2^{m_1} + 2^{m_2})$ in the 2-level hierarchical paging
 - When $m_1 = m_2 = 10$, time reduced from $O(2^{20})$ to $O(2^{11}) \Rightarrow$ 512 times less
- Extension
 - More levels in the hierarchy for larger addresses
 - Example: For 64-bit system:




2-level paging



3-level paging

Hashed Paging

- **Hash function:** Page number \rightarrow Entry in **hash page table**
 - Common in address spaces > 32 bits
 - The virtual page number is hashed into a page table
 - Entry: A linked list of elements
 - Include information for all page numbers hashable into this entry
 - Each element includes:
 - One Page number
 - Base address of corresponding frame in the memory
 - Pointer to next element in linked list



Inverted Paging

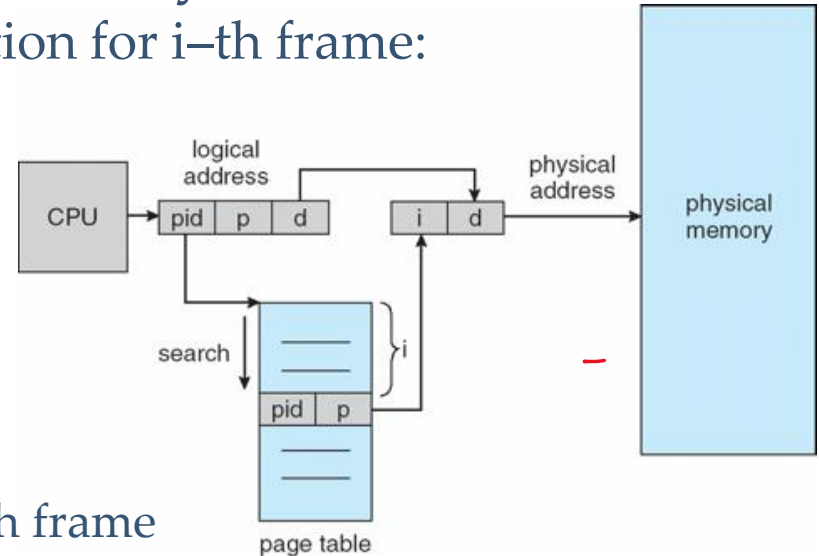
- **Inverted Page Table:** One table sorted by order of frames

- i-th entry in tables includes information for i-th frame:

- Pid of process
- Page number

- Address translation

- p extracted
- Used with pid to search table
- Match found in entry i
- Physical address = d-th address in i-th frame



- **Pro:** Decreases memory needed to store each page table

- **Con:** Increases time needed to search table

- **Possible Solutions:**

- Hash table to limit search
- TLB can accelerate access